

Proof-Carrying Review — video script

Team Two Faced Coin · IBM Bob 2.0 Hackathon · runtime 2:25 · website only, no slides · 1:45 of the solution running

TIME	ON SCREEN	VOICEOVER
0:00 – 0:12	Hero section, slow scroll	<i>Every AI code reviewer has the same problem. It gives you twenty comments and no way to tell which ones are real. So you check them yourself, and you were trying to save that time in the first place.</i>
0:12 – 0:22	Hover the hero line: “only reports one if it can write a test that fails because of it”	<i>This one has a rule. IBM Bob reports a bug only if it wrote a test that fails because of it. No failing test, no comment.</i>
0:22 – 0:40	Scroll to “Suspect. Prove. Report.” Let the four helper cards and the Prove card land.	<i>Bob reads the pull request, the ticket, the spec and the style guide. Four helpers hunt in parallel. Then every suspicion goes to a prover that has to turn it into a failing test and actually run it.</i>
0:40 – 1:05	“Pick a PR” section. Click PR 08. Let the finding card open.	<i>Here’s a real one. Pull request eight, exporting tasks to CSV. Bob says the export builds rows by joining on commas with no escaping.</i>
1:05 – 1:25	Click through to the proof test. Let the pytest output sit on screen for a beat.	<i>And here’s why you can believe it. The test Bob wrote, the real output from the run, and a task whose title has a comma in it landing in the wrong column. That’s not an opinion. It’s a receipt.</i>
1:25 – 1:40	Show the fix diff, then scroll to the dropped candidate.	<i>Bob patched the mechanical ones and reran until green. This one it threw away. It suspected quotes broke the export too, wrote the test, ran it, and the test passed. Wrong guess, dropped before anyone saw it. That’s the noise you never get.</i>
1:40 – 2:00	Scroll to the scores. Let the bars animate.	<i>Eight pull requests, sixteen bugs planted by hand, answer key kept out of the repo. Proof mode caught eighty one percent at a hundred percent precision. Zero false alarms. An ordinary single pass review caught less and handed back eight things that weren’t there.</i>
2:00 – 2:10	Obstacle course grid, sweep the green dots.	<i>Every green dot is a planted bug caught with a test attached.</i>
2:10 – 2:20	Land on the footer.	<i>It missed three, and all three were missing test cases, which a failing test can’t prove. We’d rather tell you that than hide it.</i>
2:20 – 2:25	Hold on the title.	<i>Proof Carrying Review. Code review you can check.</i>

Shooting notes

Record the click path once before the take. The PR 08 finding needs to be open and the proof-test panel rendered, or eight seconds burn on a loading state.

Pause about two seconds on the raw pytest output. It is the most convincing frame in the video and narration tends to rush past it.

Keep the browser at desktop width so the score bars sit in one row.

Running long: cut the first sentence at 0:00. Running short: after 1:40, name the baseline — a generic single-pass review, no spec, no tests.

Numbers quoted (all verifiable in the repo)

8 PRs reviewed · 16 planted bugs · 13 findings, each with a failing test · recall 81% · precision 100% · 0 false alarms
· single-pass baseline: recall 69%, precision 58%, 8 false alarms · 3 misses, all missing-test-case findings.